User Guide: Building a Reliable AI Intake Agent for the UHC Pilot

Version 1.0 – Open Narrative Intake Design

Purpose

This guide provides a practical framework for designing and piloting a **reliable AI-powered intake agent** for the first UHC sandbox test. The recommended starting point is **one customer-facing intake agent** supported by a small set of behind-the-scenes helper components for extraction, validation, summary creation, and handoff. This approach aligns with current conversational AI guidance that recommends starting with a narrowly scoped top-level workflow, building iteratively, using structured flows, and preserving strong human handoff paths rather than relying on a single unconstrained agent to do everything. [1] [2] [3] [4]

---

1. Recommended Architecture for Pilot v1

1.1 Start with one customer-facing intake agent

For the first pilot, the simplest and most reliable model is:

- **One customer-facing AI intake agent** that handles the opening conversation with the caller. [5] [6]

1.2 Use helper components behind the scenes

The customer-facing experience can be supported by a few internal workflow components:

[1] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[2] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[3] https://openai.github.io/openai-agents-python/

[4] https://developers.openai.com/api/docs/guides/agent-evals

[5] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[6] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

- **Entity extractor** to identify medications, doctors, chronic conditions, and preferences from natural speech. [7] [8]
- **Completeness checker** to determine what required information is still missing. [9] [10]
- **Confidence or risk checker** to decide whether the AI should ask another follow-up question or escalate to a human. [11] [12]
- **Summary generator** to create a concise handoff note for the agent. [13] [14] [15]
- **Routing / escalation service** to send the call to the right agent or department with the correct context. [16] [17]

1.3 Why this architecture is preferred

This pattern improves reliability because it separates the caller experience from the structured business logic, which is consistent with current guidance from Microsoft, Google,

---

[7] https://openai.github.io/openai-agents-python/

[8] https://developers.openai.com/api/docs/guides/agent-evals

[9] https://openai.github.io/openai-agents-python/

[10] https://developers.openai.com/api/docs/guides/agent-evals

[11] https://developers.openai.com/api/docs/guides/agent-evals

[12] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[13] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[14] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[15] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[16] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[17] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

and OpenAI on using tools, guardrails, workflow structure, and formal handoff context rather than one general-purpose model prompt for the entire interaction. [18] [19] [20] [21] [22]

---

2. Define the Job of the Intake Agent Clearly

2.1 Core job statement

The intake agent should have one clearly defined objective:

> **Listen to why the caller is here, identify the relevant profile details, ask only for missing critical information, and hand the conversation to a human agent with a concise summary.**

This narrow scope follows current design guidance to begin with top-level requests and avoid overloading early-stage agents with too many business tasks. [23] [24]

2.2 What the intake agent should not do in pilot v1

For the first pilot, the intake agent should **not** be expected to:

- recommend plans independently, [25] [26]

---

[18] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[19] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[20] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[21] https://openai.github.io/openai-agents-python/

[22] https://developers.openai.com/api/docs/guides/agent-evals

[23] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[24] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[25] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[26] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

- make complex eligibility or benefit determinations, [27] [28] [29]
- or serve as the primary sales or advisory experience. [30] [31]

The human agent remains responsible for the licensed, trust-based advisory part of the interaction. [32] [33]

---

3. Define the Structured Output Before Designing the Conversation

3.1 Why the schema must come first

A reliable intake workflow starts by defining exactly what the AI must produce. Without a structured output schema, the conversation may feel intelligent but still fail to generate data that is usable by the receiving agent or future systems. Structured outputs and observable workflow traces are core elements of reliable agent design and evaluation. [34] [35]

[27] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[28] https://www.cms.gov/medicare/health-drug-plans/managed-care-marketing/medicare-guidelines

[29] https://www.cms.gov/files/document/agent-broker-marketing-faqs-10-19-2022.pdf

[30] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[31] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[32] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[33] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[34] https://openai.github.io/openai-agents-python/

[35] https://developers.openai.com/api/docs/guides/agent-evals

3.2 Recommended handoff schema

The intake agent should attempt to produce a structured output that includes the following fields:

```
Caller Type: existing member | new prospect | unknown
Primary Reason for Call:
Top Priorities / Must-Haves:
Doctors:
Medications:
Chronic Conditions:
Low-Income / Medicaid History:
Anything Sensitive / Needs Clarification:
Confidence Score:
Recommended Next Step:
```

A structure like this makes it easier to evaluate completeness, compare flows, and eventually map data to CRM fields in a later phase. [36] [37]

---

4. Decide What Is Mandatory, Optional, and Out of Scope

4.1 Use the minimum necessary principle

For any workflow that may touch healthcare-related information, the design should limit collection to what is necessary for the business purpose of the pilot. HHS guidance on HIPAA's "minimum necessary" standard emphasizes limiting the use, disclosure, and request of protected health information to the minimum needed to accomplish the intended purpose. [38]

4.2 Recommended field classification for this pilot

Mandatory for handoff

The intake agent should attempt to capture:

[36] https://developers.openai.com/api/docs/guides/agent-evals

[37] https://openai.github.io/openai-agents-python/

[38] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

- reason for the call, [39] [40]
- top priorities / must-have preferences, [41] [42]
- doctors, [43] [44]
- medications, [45] [46]
- and any critical areas where the AI is uncertain. [47] [48]

Optional if naturally disclosed

The intake agent may capture:

- chronic conditions, [49] [50]
- Medicaid / low-income assistance history, [51] [52]

---

[39] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[40] https://developers.openai.com/api/docs/guides/agent-evals

[41] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[42] https://developers.openai.com/api/docs/guides/agent-evals

[43] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[44] https://developers.openai.com/api/docs/guides/agent-evals

[45] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[46] https://developers.openai.com/api/docs/guides/agent-evals

[47] https://developers.openai.com/api/docs/guides/agent-evals

[48] https://openai.github.io/openai-agents-python/

[49] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[50] https://developers.openai.com/api/docs/guides/agent-evals

[51] https://www.cms.gov/medicare/health-drug-plans/managed-care-marketing/medicare-guidelines

[52] https://www.cms.gov/files/document/agent-broker-marketing-faqs-10-19-2022.pdf

- and other context that materially affects the later conversation. [53] [54]

Out of scope for pilot v1 unless specifically required

The intake agent should not collect:

- unrelated or unnecessary details, [55]
- deeper production-sensitive information not needed in a sandbox pilot, [56] [57]
- or broader data elements that are not part of the intake objective. [58]

---

5. Design the Open Narrative Flow as a Structured State Machine

5.1 Why a structured flow matters

Even if the intake experience feels open-ended to the caller, the workflow behind it should be structured. Google recommends designing conversation agents with explicit flows and states, while Microsoft recommends clear user guidance and fallback paths that help users recover when the conversation goes off course. [59] [60]

5.2 Recommended pilot flow

The first intake style should follow this sequence:

```
State 1: Greeting + value proposition
State 2: Open narrative prompt
```

[53]https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[54]https://developers.openai.com/api/docs/guides/agent-evals

[55]https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[56]https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[57]https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[58]https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[59]https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[60]https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

```
State 3: Extract structured data
State 4: Identify missing critical fields
State 5: Ask targeted follow-up questions only for missing fields
State 6: Confirm summary
State 7: Handoff to agent with transcript + summary
```

This allows the experience to feel natural while still supporting repeatable data capture and strong evaluation. [61] [62] [63]

5.3 Example opening prompt

A recommended pilot opening is:

```
"Hi, I can gather a few details now so the agent can jump right into helping you.
In your own words, tell me why you're calling today and what matters most to you in a plan."
```

This style is consistent with guidance to explain what the AI can do, make the user benefit explicit, and create interactions that feel intuitive and human-like. [64]

---

6. Separate Conversation from Extraction

6.1 Do not rely on one step to do everything

For reliable operation, the system should separate:

1. the conversational exchange with the caller, and
2. the structured extraction of the important data elements. [65] [66]

6.2 Recommended workflow pattern

After each caller response, the system should:

---

[61] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[62] https://developers.openai.com/api/docs/guides/agent-evals

[63] https://openai.github.io/openai-agents-python/

[64] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[65] https://openai.github.io/openai-agents-python/

[66] https://developers.openai.com/api/docs/guides/agent-evals

1. store the transcript, [67] [68]
2. extract entities into the defined schema, [69] [70]
3. assign confidence levels by field, [71]
4. determine what to ask next, if anything. [72] [73]

This pattern creates a more reliable experience, especially when callers speak in long, unstructured narratives. [74] [75]

---

7. Use Field-Level Confidence Thresholds

7.1 Why confidence status matters

An intake agent should never present uncertain information as if it were fully reliable. Trust improves when the system knows what it captured confidently, what may need confirmation, and when it should escalate rather than guess. Microsoft's guidance emphasizes graceful recovery, and OpenAI's evaluation guidance highlights the value of trace-based grading to determine whether handoffs and flow decisions occurred when they should have. [76] [77]

7.2 Recommended field statuses

For each critical field, assign one of the following:

[67] https://developers.openai.com/api/docs/guides/agent-evals

[68] https://openai.github.io/openai-agents-python/

[69] https://openai.github.io/openai-agents-python/

[70] https://developers.openai.com/api/docs/guides/agent-evals

[71] https://developers.openai.com/api/docs/guides/agent-evals

[72] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[73] https://developers.openai.com/api/docs/guides/agent-evals

[74] https://openai.github.io/openai-agents-python/

[75] https://developers.openai.com/api/docs/guides/agent-evals

[76] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[77] https://developers.openai.com/api/docs/guides/agent-evals

- **Captured** [78]
- **Likely captured / needs confirmation** [79]
- **Missing** [80]
- **Unknown / do not know** [81]

7.3 Recommended decision rules

The workflow should support rules such as:

```
If medications confidence is below threshold:
    ask a follow-up question
If caller shows frustration:
    stop follow-ups and hand off
If summary quality is low:
    flag "needs verification" for the human agent
```

This approach reduces false confidence and gives agents better visibility into what still needs validation. [82] [83] [84]

---

8. Create Purpose-Built Follow-Up Micro-Scripts

8.1 Why follow-up quality matters

The success of the open narrative intake style depends less on the first question than on the quality of the system's follow-up questions. Good follow-ups should be short, specific, and clearly tied to missing information. Microsoft's conversation design guidance emphasizes reducing cognitive load and guiding users efficiently through the experience. [85]

---

[78] https://developers.openai.com/api/docs/guides/agent-evals

[79] https://developers.openai.com/api/docs/guides/agent-evals

[80] https://developers.openai.com/api/docs/guides/agent-evals

[81] https://developers.openai.com/api/docs/guides/agent-evals

[82] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[83] https://developers.openai.com/api/docs/guides/agent-evals

[84] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[85] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

8.2 Example follow-up scripts

If no must-have preference was captured

```
"What matters most to you in a plan – keeping certain doctors, lower prescription costs, or something else?"
```

[86]

If no doctors were identified

```
"Who are the main doctors you want to keep seeing?"
```

[87]

If no medications were identified

```
"What prescriptions do you take regularly?"
```

[88]

If chronic-condition context is unclear

```
"Are there any ongoing health conditions that should be considered when looking at coverage?"
```

[89]

If affordability or assistance history needs to be clarified

```
"Have you ever received Medicaid or any low-income assistance that we should take into account?"
```

[90] [91]

---

[86] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[87] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[88] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[89] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[90] https://www.cms.gov/medicare/health-drug-plans/managed-care-marketing/medicare-guidelines

[91] https://www.cms.gov/files/document/agent-broker-marketing-faqs-10-19-2022.pdf

These prompts should remain voice-friendly, easy to answer, and limited to missing business-critical information. [92] [93]

---

9. Make Summary Confirmation a Required Reliability Step in Pilot v1

9.1 Why summary confirmation should be on

For the first pilot, summary confirmation should be enabled because it improves transparency, helps callers correct errors, and increases the likelihood that the agent receives a usable handoff. Microsoft recommends clearly explaining AI behavior and limitations, while OpenAI recommends structured evaluation to compare whether changes such as confirmation meaningfully reduce error or rework. [94] [95]

9.2 Recommended summary confirmation pattern

```
“Here’s what I captured: you’re looking for a plan that keeps Dr. Patel in-network,
covers two regular prescriptions, and keeps costs predictable. I also noted that you’ve had
low-income assistance in the past. Did I get that right?”
```

This creates a trust moment for the caller, improves data accuracy, and strengthens the agent handoff. [96] [97]

---

[92] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[93] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[94] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[95] https://developers.openai.com/api/docs/guides/agent-evals

[96] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[97] https://developers.openai.com/api/docs/guides/agent-evals

10. Treat the Handoff as a Product Requirement

10.1 Why handoff quality matters

Microsoft's handoff guidance specifies that a bot should pass both context and transcript information so the human agent can continue the conversation without losing what has already happened. Recent Twilio guidance similarly emphasizes that a seamless handoff should preserve context and provide a summary so the customer does not need to repeat the problem. [98] [99] [100]

10.2 Recommended handoff payload

The handoff package should include:

- structured summary, [101] [102] [103]
- transcript or transcript excerpt, [104] [105]
- confidence or "needs verification" flags, [106] [107]

---

[98] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[99] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[100] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[101] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[102] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[103] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[104] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[105] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[106] https://developers.openai.com/api/docs/guides/agent-evals

[107] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

- routing recommendation, [108]
- elapsed intake time, [109]
- and unresolved questions. [110] [111]

10.3 Design for 5-second readability

The receiving agent should see a short skim-friendly summary, not a long narrative paragraph. A concise structured card is more likely to be used in real time and is more consistent with effective handoff design. [112] [113] [114]

Example agent card

```
Reason: Shopping for plan / wants predictable cost
Priorities: keep PCP, manage Rx costs
Doctors: Dr. Patel, Dr. Singh
Meds: Eliquis, Metformin
Potential constraints: prior low-income assistance
Needs verification: specialist list
```

11. Add Privacy, Verification, and Escalation Guardrails

11.1 Privacy and minimum necessary

If member-related information is involved, the pilot should include policies and controls that align with minimum necessary handling of health information. HHS guidance states that

[108] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[109] https://developers.openai.com/api/docs/guides/agent-evals

[110] https://developers.openai.com/api/docs/guides/agent-evals

[111] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[112] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[113] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[114] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

covered entities should design policies and procedures that limit access and disclosure to what is needed for the intended function. [115]

11.2 Medicare-related operational awareness

If future versions of this workflow touch Medicare marketing, sales, or enrollment processes, compliance and legal review should confirm how CMS communication and recording requirements apply before any broader deployment. CMS publishes Medicare marketing guidance, and CMS FAQs describe recording requirements that apply to TPMO-related beneficiary calls in certain contexts. [116] [117]

11.3 Recommended guardrails for the pilot

Pilot v1 should include:

- sandboxed or simulated data wherever possible, [118] [119]
- explicit rules for when identity verification is required before referencing known member information, [120]
- clear escalation language when handing off to a human, [121] [122]

[115] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[116] https://www.cms.gov/medicare/health-drug-plans/managed-care-marketing/medicare-guidelines

[117] https://www.cms.gov/files/document/agent-broker-marketing-faqs-10-19-2022.pdf

[118] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[119] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[120] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[121] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[122] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

- and an audit trail of what was captured, confirmed, and transferred. [123] [124]

---

12. Build an Evaluation Dataset Before Launch

12.1 Reliability must be demonstrated, not assumed

A reliable intake agent is one that performs consistently across a range of realistic caller behaviors. OpenAI recommends moving from individual traces to structured datasets and evaluation runs, while Google recommends iterative testing with formal test cases as the agent design evolves. [125] [126]

12.2 Recommended pilot test set

Before launching to the full 30–40 participant pilot, create a dataset of at least **25–40 realistic scenarios** that includes:

- concise callers, [127] [128]
- rambling callers, [129] [130]
- callers who shift topics, [131] [132]
- unclear medication names, [133] [134]
- callers who do not know key details, [135]

---

[123] https://developers.openai.com/api/docs/guides/agent-evals

[124] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[125] https://developers.openai.com/api/docs/guides/agent-evals

[126] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[127] https://developers.openai.com/api/docs/guides/agent-evals

[128] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[129] https://developers.openai.com/api/docs/guides/agent-evals

[130] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[131] https://developers.openai.com/api/docs/guides/agent-evals

[132] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[133] https://developers.openai.com/api/docs/guides/agent-evals

[134] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[135] https://developers.openai.com/api/docs/guides/agent-evals

- frustrated callers, [136] [137]
- existing members versus new prospects, [138] [139]
- and affordability-sensitive scenarios. [140] [141]

12.3 Recommended grading questions

Each run should be scored against consistent criteria such as:

- Did the AI capture the reason for the call? [142]
- Did it capture doctors and medications? [143]
- Did it ask only needed follow-up questions? [144] [145]
- Did it avoid unnecessary questions? [146] [147]
- Did it summarize accurately? [148]

---

[136] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[137] https://developers.openai.com/api/docs/guides/agent-evals

[138] https://developers.openai.com/api/docs/guides/agent-evals

[139] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[140] https://www.cms.gov/medicare/health-drug-plans/managed-care-marketing/medicare-guidelines

[141] https://www.cms.gov/files/document/agent-broker-marketing-faqs-10-19-2022.pdf

[142] https://developers.openai.com/api/docs/guides/agent-evals

[143] https://developers.openai.com/api/docs/guides/agent-evals

[144] https://developers.openai.com/api/docs/guides/agent-evals

[145] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[146] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[147] https://developers.openai.com/api/docs/guides/agent-evals

[148] https://developers.openai.com/api/docs/guides/agent-evals

- Did it hand off when uncertain? [149] [150]
- Did the human agent still need to repeat the full intake? [151] [152] [153]

---

13. Instrument the Workflow with Traces and Failure Tags

13.1 Why instrumentation is essential

Workflow traces make it easier to identify where the experience breaks down, such as missed extraction, poor follow-up logic, or weak handoff quality. OpenAI recommends using traces first for workflow debugging and then graduating to more structured evaluators as the workflow matures. [154] [155]

13.2 Recommended failure tags

The team should tag failures consistently using labels such as:

- missed_doctor_capture [156]
- missed_med_capture [157]
- summary_wrong [158]

[149] https://developers.openai.com/api/docs/guides/agent-evals

[150] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[151] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[152] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[153] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[154] https://developers.openai.com/api/docs/guides/agent-evals

[155] https://openai.github.io/openai-agents-python/

[156] https://developers.openai.com/api/docs/guides/agent-evals

[157] https://developers.openai.com/api/docs/guides/agent-evals

[158] https://developers.openai.com/api/docs/guides/agent-evals

- followup_too_many [159] [160]
- caller_confused [161]
- caller_frustrated [162] [163]
- agent_reasked_everything [164] [165] [166]
- handoff_missing_context [167]

These labels allow the team to improve based on specific failure patterns rather than general feedback like "it felt choppy." [168] [169]

---

14. Keep the First Version Simple

14.1 Do not start with multiple visible agents

Although modern agent frameworks support multi-agent orchestration and handoff between specialist agents, the first pilot should avoid exposing that complexity to callers unless a clear business need has already been established. Current design guidance

[159] https://developers.openai.com/api/docs/guides/agent-evals

[160] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[161] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[162] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[163] https://developers.openai.com/api/docs/guides/agent-evals

[164] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[165] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[166] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[167] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[168] https://developers.openai.com/api/docs/guides/agent-evals

[169] https://openai.github.io/openai-agents-python/

recommends starting with a manageable workflow and expanding only after the top-level path has been validated. [170] [171] [172]

14.2 Recommended pilot v1 model

Pilot v1 should include:

- **one caller-facing intake agent**, [173] [174]
- **three invisible helper functions** for extraction, completeness/confidence checking, and summary generation, [175] [176]
- and strict escalation rules when confidence is low or the caller is frustrated. [177] [178]

14.3 When to consider expansion

Additional specialist agents should be considered only after the pilot proves that the intake workflow itself is useful and that separate domains truly require distinct logic or policy controls. [179] [180]

---

[170] https://openai.github.io/openai-agents-python/

[171] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[172] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[173] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[174] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[175] https://openai.github.io/openai-agents-python/

[176] https://developers.openai.com/api/docs/guides/agent-evals

[177] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[178] https://developers.openai.com/api/docs/guides/agent-evals

[179] https://openai.github.io/openai-agents-python/

[180] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

15. Recommended Build Sequence

Phase 1 — Define

1. Lock the exact purpose of the intake agent. [181] [182]
2. Define the structured output schema. [183] [184]
3. Classify fields as mandatory, optional, or out of scope. [185] [186]
4. Define the agent handoff payload. [187] [188] [189]

Phase 2 — Design

1. Write the state machine for the open narrative flow. [190] [191]
2. Draft the follow-up micro-scripts. [192]

---

[181] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[182] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[183] https://openai.github.io/openai-agents-python/

[184] https://developers.openai.com/api/docs/guides/agent-evals

[185] https://www.hhs.gov/hipaa/for-professionals/privacy/guidance/minimum-necessary-requirement/index.html

[186] https://developers.openai.com/api/docs/guides/agent-evals

[187] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[188] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[189] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[190] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[191] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[192] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

3. Draft the summary confirmation script. [193] [194]
4. Draft the human agent warm-handoff script. [195] [196] [197]

Phase 3 — Build

1. Build the conversational intake agent. [198] [199]
2. Build the extraction layer. [200] [201]
3. Add confidence thresholds and follow-up logic. [202] [203]
4. Add summary generation and routing. [204] [205] [206]

---

193 https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

194 https://developers.openai.com/api/docs/guides/agent-evals

195 https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

196 https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

197 https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

198 https://openai.github.io/openai-agents-python/

199 https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

200 https://openai.github.io/openai-agents-python/

201 https://developers.openai.com/api/docs/guides/agent-evals

202 https://developers.openai.com/api/docs/guides/agent-evals

203 https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

204 https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

205 https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

206 https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

5. Add logging, tracing, and transcript capture. [207] [208]

Phase 4 — Validate

1. Run internal dry-run scenarios. [209] [210]
2. Measure capture accuracy and missing-field rates. [211]
3. Tune prompts, micro-scripts, and thresholds. [212] [213]
4. Retest the same scenarios to confirm improvement. [214] [215]

Phase 5 — Pilot

1. Run the 30–40 participant pilot. [216] [217]
2. Review caller feedback, agent feedback, and trace data together. [218] [219]
3. Decide whether to keep, refine, or redesign the intake style based on evidence. [220] [221]

---

16. Short Summary: What Makes the Intake Agent Reliable?

A reliable intake agent is **not just a good prompt**. It is:

---

[207] https://developers.openai.com/api/docs/guides/agent-evals

[208] https://openai.github.io/openai-agents-python/

[209] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[210] https://developers.openai.com/api/docs/guides/agent-evals

[211] https://developers.openai.com/api/docs/guides/agent-evals

[212] https://developers.openai.com/api/docs/guides/agent-evals

[213] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[214] https://developers.openai.com/api/docs/guides/agent-evals

[215] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[216] https://developers.openai.com/api/docs/guides/agent-evals

[217] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[218] https://developers.openai.com/api/docs/guides/agent-evals

[219] https://openai.github.io/openai-agents-python/

[220] https://developers.openai.com/api/docs/guides/agent-evals

[221] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

- a tightly scoped workflow, [222] [223]
- with a defined output schema, [224] [225]
- targeted follow-up logic, [226] [227]
- field-level confidence thresholds, [228]
- summary confirmation, [229] [230]
- strong handoff context, [231] [232] [233]
- and a formal evaluation loop using traces and test cases. [234] [235] [236]

---

[222] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[223] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[224] https://openai.github.io/openai-agents-python/

[225] https://developers.openai.com/api/docs/guides/agent-evals

[226] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[227] https://developers.openai.com/api/docs/guides/agent-evals

[228] https://developers.openai.com/api/docs/guides/agent-evals

[229] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[230] https://developers.openai.com/api/docs/guides/agent-evals

[231] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[232] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[233] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[234] https://developers.openai.com/api/docs/guides/agent-evals

[235] https://openai.github.io/openai-agents-python/

[236] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

17. Recommended Pilot v1 Configuration for UHC

For the first UHC sandbox pilot, the recommended configuration is:

Caller-facing experience

- One AI intake agent using the **open narrative** opening. [237] [238]
- Follow-up questions only when critical information is missing. [239] [240]
- Summary confirmation enabled. [241] [242]
- Warm handoff to the human agent. [243] [244] [245]

Behind the scenes

- One extractor, [246] [247]
- one completeness/confidence checker, [248]

[237] https://docs.cloud.google.com/dialogflow/cx/docs/concept/agent-design

[238] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[239] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[240] https://developers.openai.com/api/docs/guides/agent-evals

[241] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[242] https://developers.openai.com/api/docs/guides/agent-evals

[243] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[244] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[245] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[246] https://openai.github.io/openai-agents-python/

[247] https://developers.openai.com/api/docs/guides/agent-evals

[248] https://developers.openai.com/api/docs/guides/agent-evals

- and one summary generator. [249] [250] [251]

Hard operating rules

- If confidence is low, ask one targeted follow-up. [252] [253]
- If confidence remains low or the caller becomes frustrated, hand off immediately with a “needs verification” flag to the agent. [254] [255] [256]

If you’d like, I can take this one step further and turn it into either:

1. a **Word-ready internal SOP / playbook**,
2. a **one-page executive summary**, or
3. a **RACI + implementation checklist** for Jared / Nicole / TP / UHC stakeholders.

---

[249] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0

[250] https://www.twilio.com/en-us/blog/developers/best-practices/ai-human-handoff-tac-orchestrator-studio-flex

[251] https://www.twilio.com/en-us/blog/developers/tutorials/product/twilio-agent-connect-flex-human-handoff

[252] https://developers.openai.com/api/docs/guides/agent-evals

[253] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[254] https://learn.microsoft.com/en-us/power-platform/well-architected/experience-optimization/conversation-design

[255] https://developers.openai.com/api/docs/guides/agent-evals

[256] https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-pattern-handoff-human?view=azure-bot-service-4.0